

# 1 2019-01-15-applications

Let  $z_1, \dots, z_n \in \mathbb{Z}_+$  be a set of  $n$  non-negative integers, i.e.  $\mathbb{Z}_+ = \{0, 1, 2, \dots\}$ . Define for any positive  $\mu > 0$  the function

$$f(\mu) = \sum_{i=1}^n \mu - z_i \log \mu,$$

where  $\sum_{i=1}^n$  denotes the sum of  $n$  terms, and  $\log$  is the natural logarithm ( $\log = \ln$ ).

1. Derive an expression in terms of  $\mu, n, z_i$  for the first derivative  $f'(\mu)$ .
2. Derive an expression in terms of  $n, z_i$  for the value of  $\mu$  which is a critical point of  $f$ . Hint: set  $f'(\mu) = 0$  and solve for  $\mu$ .
3. If  $n = 3$  and  $z_1 = 1, z_2 = 0, z_3 = 2$ , then compute the value of the critical point  $\mu = \underline{\hspace{1cm}}$  and the critical function value  $f(\mu) = \underline{\hspace{1cm}}$ . Fill in the blanks here but show your work below.
4. Is the critical point a minimum ( $f''(\mu) > 0$ ), maximum ( $f''(\mu) < 0$ ), or inflection point ( $f''(\mu) = 0$ )?
5. Write a function in the C programming language that computes the critical point  $\mu$ . The function should have two input arguments: the number  $n$  of data `int n` and a pointer `int* z` to the data  $z_i$ . The function should output the critical point  $\mu$  as a `double`.

## 2 2019-01-17-nearest-neighbors

Here are  $n = 5$  data with  $p = 2$  input dimensions. Each row is a person for which we have measured the height (first column of  $X$ , in centimeters), weight (second column in pounds). The output  $y$  that we want to predict is diabetes diagnosis status (1=diabetes, 0=not).

$$X = \begin{bmatrix} 170 & 140 \\ 200 & 160 \\ 180 & 200 \\ 140 & 150 \\ 150 & 130 \end{bmatrix}, \quad y = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 1 \end{bmatrix}, \quad d = \begin{bmatrix} \_\_\_\_\_\_ \\ \_\_\_\_\_\_ \\ \_\_\_\_\_\_ \\ \_\_\_\_\_\_ \\ \_\_\_\_\_\_ \end{bmatrix}$$

The goal is to compute the class predicted by the  $K = 3$  nearest neighbor model, for a new/test person with features  $x$ =[height = 160 cm, weight = 130 pounds].

Assume that we use the Manhattan/L1 distance metric,

$$d(x, x') = \sum_{j=1}^p |x_j - x'_j|,$$

i.e. the total distance between  $x, x'$  is the sum of distances on each of the two component dimensions (height and weight).

1. Fill in the blanks in the vector of distances  $d$  above (each row should be the Manhattan/L1 distance between the new/test person and the training data).
2. Which are the three nearest neighbors? (write a star\* for each of the three nearest neighbors in the blank next to the corresponding distances/rows)
3. What is the overall predicted class? (0 or 1)

### 3 2019-01-22-pseudocode

Below I have written pseudo-code for a version of the  $k$ -nearest neighbors algorithm. Fill in the blank on line 12 so that the algorithm computes predictions  $\hat{y}_k$  for all  $k \in \{1, \dots, K_{\max}\}$ .

```
1: Function PRED1TOKMAXNEARESTNEIGHBORS
2: Inputs: train inputs/features  $x_1, \dots, x_n$ , outputs/labels  $y_1, \dots, y_n$ ,
3:   test input/feature  $x'$ , max number of neighbors  $K_{\max}$ :
4: for  $i = 1$  to  $n$  do:
5:    $d_i \leftarrow \text{DISTANCE}(x', x_i)$ 
6: end for
7:  $t_1, \dots, t_n \leftarrow \text{SORTEDINDICES}(d_1, \dots, d_n)$ 
8: totalY  $\leftarrow 0.0$ 
9: for  $k = 1$  to  $K_{\max}$  do:
10:    $i \leftarrow t_k$ 
11:   totalY  $\mathrel{+}= y_i$ 
12:    $\hat{y}_k \leftarrow \underline{\hspace{2cm}}$ 
13: end for
14: Output: predictions  $\hat{y}_1, \dots, \hat{y}_{K_{\max}}$ .
```

## 4 2019-01-24-cross-validation

The image below represents a training data set with  $n = 70$  observations, one for each individual image of a digit. In order to perform cross-validation, fold ID numbers  $\in \{1, 2, 3\}$  have been assigned to all observations/images in the corresponding row/letter.



1. For fold/split 1 which observations/letters are used for the training set? \_\_\_\_\_  
Which observations/letters are used for validation set? \_\_\_\_\_
2. For fold/split 2. Training set = \_\_\_\_\_, Validation set = \_\_\_\_\_.
3. For fold/split 3. Training set = \_\_\_\_\_, Validation set = \_\_\_\_\_.

## 5 2019-01-29-nearest-neighbors-code

Here are  $n = 4$  data with  $p = 2$  input dimensions. Each row is a person for which we have measured the height (first column of  $X$ , in centimeters), weight (second column in pounds). The output  $y$  that we want to predict is a blood pressure measurement.

$$X = \begin{bmatrix} 170 & 140 \\ 200 & 160 \\ 180 & 200 \\ 140 & 150 \end{bmatrix}, y = \begin{bmatrix} 120 \\ 115 \\ 135 \\ 140 \end{bmatrix}$$

1. How would you represent these data in R? (fill in the blanks)

```
X <- matrix(c(____,____,____,____,____,____,____,____),  
nrow=____, ncol=____)  
y <- c(____,____,____,____)
```

2. Using `.C("myPrint_interface", as.double(X), PACKAGE="myPkg")` we can access the inputs via a C++ function:

```
void myPrint_interface(double *X_ptr){  
    ...  
}
```

Inside that function, what is the value of `X_ptr[4]`? \_\_\_\_\_

## 6 2019-01-31-coding

Here is a block of C++ code which declares some variables that will be used for computing the nearest neighbors predictions for a multi-class classification problem with `n_labels=10` classes. Assume there are `nrow=1000` training observations in a feature/input space of size `ncol=256`. For each line of code, indicate (1) the total number of elements stored in the corresponding C array, (2) the underlying C type of each of those elements, double or int, and (3) YES if that line of code performs a dynamic memory allocation to get a new C array, otherwise NO.

```
Eigen::Map< Eigen::MatrixXd > train_inputs_mat(train_inputs_ptr, nrow, ncol);
```

(1)\_\_\_\_\_ (2)\_\_\_\_\_ (3)\_\_\_\_\_

```
Eigen::Map< Eigen::VectorXd > test_input_vec(test_input_ptr, ncol);
```

(1)\_\_\_\_\_ (2)\_\_\_\_\_ (3)\_\_\_\_\_

```
Eigen::VectorXd distance_vec(nrow);
```

(1)\_\_\_\_\_ (2)\_\_\_\_\_ (3)\_\_\_\_\_

```
Eigen::VectorXi sorted_index_vec(nrow);
```

(1)\_\_\_\_\_ (2)\_\_\_\_\_ (3)\_\_\_\_\_

```
Eigen::VectorXi label_count_vec(n_labels);
```

(1)\_\_\_\_\_ (2)\_\_\_\_\_ (3)\_\_\_\_\_

## 7 2019-02-05-linear-regression

Let  $w \in \mathbb{R}$  and  $g(w) = \frac{1}{2}(w - 4)^2$  be a cost function that we will minimize via gradient descent. Derive an expression for gradient  $\nabla g(w)$  in terms of  $w$ .

If we start at the origin  $w^{(0)} = 0$ , what is the starting value of the cost function?  $g(w^{(0)}) = \underline{\hspace{2cm}}$

Let's use gradient descent with step size  $\alpha = 0.5$  to find a lower cost. Recall that the gradient descent updates are given by  $w^{(t)} = w^{(t-1)} - \alpha \nabla g(w^{(t-1)})$  for all  $t \in \{1, 2, \dots\}$ . What are the first two steps in gradient descent? Fill in the blanks below. (each should be a real number)

$\nabla g(w^{(0)}) = \underline{\hspace{2cm}}$

$w^{(1)} = \underline{\hspace{2cm}}$

$\nabla g(w^{(1)}) = \underline{\hspace{2cm}}$

$w^{(2)} = \underline{\hspace{2cm}}$

What is the ending value of the cost function?  $g(w^{(2)}) = \underline{\hspace{2cm}}$

## 8 2019-02-07-logistic-regression

**Poisson regression** is a machine learning problem where the output/label  $y_i \in \{0, 1, \dots\}$  is integer-valued, and the input/features  $x_i \in \mathbb{R}^p$  is a real vector as usual. For example  $y_i$  could be the number of pennies in your wallet, the number of cars in your garage, or the number of books in your backpack — all of these are non-negative integers.

This case needs special treatment because if you use standard linear regression, with the square loss, you end up with a prediction function  $f(x_i) \in \mathbb{R}$  that predicts real numbers, and it does not make sense to predict a negative number (or a non-integer number) of pennies/cars/books. We will derive a loss function to use in this case.

We assume  $y_i \sim \text{Poisson}(\lambda_i)$  where  $\lambda_i \in \mathbb{R}_+$  is a non-negative real number — it is called the mean or rate parameter. The Poisson probability mass function is

$$\Pr(y_i, \lambda_i) = \frac{\lambda_i^{y_i} e^{-\lambda_i}}{y_i!}$$

Derive an expression in terms of  $y_i$  and  $\lambda_i$  for the log-likelihood of the mean parameter  $\lambda_i$  given a single label  $y_i$ :

$$\log \Pr(y_i, \lambda_i) = \underline{\hspace{15cm}}$$

We learn a linear function  $f(x_i) = w^T x_i = \log \lambda_i \in \mathbb{R}$ , which means that  $\lambda_i = e^{w^T x_i}$ .

The negative log-likelihood of a particular weight vector  $w \in \mathbb{R}^p$  is therefore

$$\begin{aligned} -\text{LogLik}(w) &= -\sum_{i=1}^n \log \Pr(y_i, e^{w^T x_i}) \\ &= \sum_{i=1}^n \underline{\hspace{15cm}} \end{aligned}$$

In the blank above, write an expression for the negative log-likelihood in terms of  $y_i, x_i, w$ . There should be three terms that are added/subtracted together. Circle the term that does NOT depend on  $w$  — the other two terms can be used as a loss function to minimize.



## 9 2019-02-12-log-reg-gradient

**Logistic regression** is a machine learning problem where the output/label  $\tilde{y}_i \in \{-1, 1\}$  is binary-valued, and the inputs/features  $x_i \in \mathbb{R}^p$  is a real vector as usual.

Let  $X \in \mathbb{R}^{n \times p}$  be the input/feature matrix, and let  $\tilde{y} \in \{-1, 1\}^n$  be the vector of labels. Let

$$\tilde{Y} = \text{Diag}(\tilde{y}) = \begin{bmatrix} \tilde{y}_1 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \tilde{y}_n \end{bmatrix} \in \mathbb{R}^{n \times n}$$

be a matrix with labels on the diagonal and zeros elsewhere.

The total logistic loss for the linear prediction function  $f(x_i) = w^T x_i$  is

$$\mathcal{L}(w) = \sum_{i=1}^n \log[1 + \exp(-\tilde{y}_i w^T x_i)].$$

Let

$$v = \begin{bmatrix} \frac{1}{1 + \exp(\tilde{y}_1 w^T x_1)} \\ \vdots \\ \frac{1}{1 + \exp(\tilde{y}_n w^T x_n)} \end{bmatrix} \in \mathbb{R}^n.$$

Derive an expression in terms of  $X, \tilde{Y}, v$  for the gradient of the total logistic loss and put it in the blank below.

$$\nabla \mathcal{L}(w) = \underline{\hspace{10cm}}$$

## 10 2019-02-14-L2-regularization

In the statistics literature, the ridge regression problem is typically defined as follows. The output/label  $y_i \in \mathbb{R}$  is real-valued, and the inputs/features  $x_i \in \mathbb{R}^p$  is a real vector as usual. Let  $X \in \mathbb{R}^{n \times p}$  be the input/feature matrix, and let  $y \in \mathbb{R}^n$  be the vector of labels.

The linear prediction function is  $f_{\beta,w}(x_i) = \beta + w^T x_i$ , where  $\beta \in \mathbb{R}$  is called the “intercept” or “bias” term, and  $w \in \mathbb{R}^p$  is the usual vector of weights, one for each feature.

Let  $\mathbf{1}_n = [1 \ \cdots \ 1]^T \in \mathbb{R}^n$  be an  $n$ -vector of ones. The ridge regression cost function can then be defined as

$$\mathcal{C}_\lambda(\beta, w) = \|\mathbf{1}_n \beta + Xw - y\|_2^2 + \lambda \|w\|_2^2.$$

Note in the definition above that L2 regularization is only used for the weight vector  $w$  (not for the bias/intercept  $\beta$ ).

The optimal model parameters for a particular  $\lambda \geq 0$  are defined as

$$\hat{\beta}^\lambda, \hat{w}^\lambda = \arg \min_{\beta \in \mathbb{R}, w \in \mathbb{R}^p} \mathcal{C}_\lambda(\beta, w).$$

To find the optimal model parameters we must first compute the gradients, (fill in the blanks below in terms of  $X, y, w, \beta, \mathbf{1}_n$ )

$$\nabla_\beta \mathcal{C}_\lambda(\beta, w) = \underline{\hspace{4cm}}$$

$$\nabla_w \mathcal{C}_\lambda(\beta, w) = \underline{\hspace{4cm}}$$

## 11 2019-02-26-line-search

Exact line search in 2 dimensions. For  $w \in \mathbb{R}^2$ , define the cost function

$$C(w) = \frac{1}{2}(w_1 - 1)^2 + \frac{1}{2}(w_2 + 1)^2 = \frac{1}{2} \|w + \begin{bmatrix} -1 \\ 1 \end{bmatrix}\|_2^2$$

Derive an expression for the gradient in terms of  $w$ ,  $\nabla C(w) =$  \_\_\_\_\_  
Let  $w^{(0)} = 0$  be the starting point of gradient descent, at the origin.

The descent direction is  
 $d^{(0)} = -\nabla C(w^{(0)}) =$  \_\_\_\_\_  
The cost of a step with size  $\alpha > 0$  in that direction is

$$\mathcal{C}_0(\alpha) = C(w^{(0)} + \alpha d^{(0)}).$$

To find the step size with the lowest cost we first need the derivative (in terms of  $\alpha$ ):

$$\mathcal{C}'_0(\alpha) =$$

Setting the derivative to zero,  $\mathcal{C}'_0(\alpha) = 0$ , then solving  
for  $\alpha$  implies an optimal step size of  $\alpha^{(0)} = \arg \min_{\alpha} \mathcal{C}_0(\alpha) =$  \_\_\_\_\_

Taking that step lands us at

$$w^{(1)} = w^{(0)} + \alpha^{(0)} d^{(0)} =$$

which has a cost of

$$\mathcal{C}_0(\alpha^{(0)}) = C(w^{(1)}) =$$

## 12 2019-02-28-backtracking-line-search

**Exact line search for un-regularized least squares linear regression.**

For an input/feature matrix  $X \in \mathbb{R}^{n \times p}$ , an output/label vector  $y \in \mathbb{R}^n$ , and a weight vector  $w \in \mathbb{R}^p$ , define the least squares loss function

$$L(w) = \frac{1}{2} \|Xw - y\|_2^2$$

Derive an expression for the gradient in terms of  $X, y, w$ .  $\nabla L(w) =$ \_\_\_\_\_

The descent direction at  $w$  is the negative gradient,  $d = -\nabla L(w) =$ \_\_\_\_\_  
The cost of a step with size  $\alpha > 0$  in that direction is

$$\mathcal{L}(\alpha) = L(w + \alpha d) =$$

To find the step size with the min cost we first need the derivative (in terms of  $\alpha, d, w, X, y$ ):

$$\mathcal{L}'(\alpha) =$$

Setting the derivative to zero,  $\mathcal{L}'(\alpha) = 0$ , implies an optimal step size of

$$\arg \min_{\alpha} \mathcal{L}'(\alpha) =$$